

Ex. No: 1	
16th July, 2024	

Linux Commands

Aim:

To find and implement basic Linux Commands.

Programs and Outputs:

1. Display information about files in the current directory.

```
anantha@LAPTOP
$ ls
cat
```

2. Display the current working directory.

```
/cat$ pwd
/mnt/d/SAM/Sem-7/Big D/cat
```

3. Create a new directory.

```
anantha@LAPTOP
$ mkdir cat
anantha@LAPTOP
$ ls
cat
```

4. Navigate between different folders.

```
anantha@LAPTOP
$ cd cat
```

5. Remove empty directories from the directory lists.

```
anantha@LAPTOP
$ ls
cat  dog
anantha@LAPTOP
$ rmdir dog
anantha@LAPTOP
$ ls
cat
```

6. a. Copy files from one directory to the same directory.

```
/cat$ cat pink.txt
Do pink cats even exist?
/cat$ cp orange.txt pink.txt
/cat$ cat pink.txt
Orange cats are not normal
```

6. b. Copy files from one directory to another directory.

```
/cow$ ls  
orange.txt  
/cat$ cp pink.txt /mnt/d/SAM/Sem-7/'Big D'/cow  
/cow$ ls  
orange.txt pink.txt
```

7. a. Rename a filename to another name.

```
mv orange.txt or.txt  
anantha@LAPTOP-M  
ls  
or.txt pink.txt
```

7. b. Move a file from one directory to another.

```
cow$ ls  
orange.txt pink.txt  
cat$ mv or.txt /mnt/d/SAM/Sem-7/'Big D'/cow  
cow$ ls  
or.txt orange.txt pink.txt
```

8. a. Delete individual files from a directory

```
cow$ ls  
or.txt orange.txt pink.txt  
anantha@LAPTOP  
cow$ rm or.txt  
cow$ ls  
orange.txt pink.txt
```

8. b. Delete an entire directory which contains a few files.

```
anantha@LAPTOP  
em-7/Big D$ ls  
cat cow  
rm -r cow  
/Sem-7/Big D$ ls  
cat
```

9. Get basic information about the OS

```
em-7/Big D$ uname -a  
Microsoft-WSL2 #1 SMP Fri Mar  
29 23:14:13 UTC 2024 x86_64 x86_64 x  
86_64 GNU/Linux
```

10. Find a file in the directory.

```
7/Big D$ find /mnt/d/SAM/Sem-7/'Big D'/  
cat -name pink.txt  
/mnt/d/SAM/Sem-7/Big D/cat/pink.txt
```

11. Create empty files

```
/cat$ touch orange.txt
```

12. Display file contents on terminal

```
/cat$ cat orange.txt  
Orange cats are not normal
```

13. Clear terminal

```
Sem-7/Big D$ clear
```

14. Display the processes in terminal

```
/cow$ ps
    PID TTY          TIME CMD
      361 pts/0    00:00:00 bash
     7409 pts/0    00:00:00 ps
```

15. Access manual for all Linux commands

```
/cow$ man cd
```

cd(n)	Tcl Built-In Commands	cd(n)
<hr/>		
NAME	cd - Change working directory	
SYNOPSIS	cd ? <u>dirName</u> ?	
<hr/>		
DESCRIPTION	Change the current working directory to <u>dirName</u> , or to the home directory (as specified in the HOME environment variable) if <u>dirName</u> is not given. Returns an empty string. Note that the current working directory is a per-process resource; the cd command changes the working directory for all interpreters and (in a threaded environment) all threads.	
Manual page cd(n) line 1 (press h for help or q to quit)		

16. Search for a specific string in an output

```
cat$ grep cats pink.txt
Orange cats are not normal
```

17. Display active processes on the terminal

```
/cow$ top
```

top - 10:54:44 up 24 min, 1 user, load average: 0.09, 0.08, 0.01	Tasks: 42 total, 1 running, 41 sleeping, 0 stopped, 0 zombie	%Cpu(s): 0.1 us, 0.0 sy, 0.0 ni, 99.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st	MiB Mem : 7801.5 total, 6992.2 free, 569.8 used, 239.5 buff/cache	MiB Swap: 2048.0 total, 2048.0 free, 0.0 used. 7000.1 avail Mem							
<hr/>											
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	167168	12376	8024	S	1.0	0.2	0:15.94	systemd
top - 10:54:50 up 24 min, 1 user, load aver											
Tasks: 34 total, 1 running, 33 sleeping,											

18. Download files from the internet.

```
cat$ wget https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/
```

Resolving archive.ics.uci.edu (archive.ics.uci.edu)... 128.195.10.252
Connecting to archive.ics.uci.edu (archive.ics.uci.edu)|128.195.10.252|:443... connected.

19. Create or update passwords for existing users

```
/cat$ passwd anantha
```

Changing password for anantha.
Current password:

20. View the exact location of any tool/software installed

```
anantha@LAPTOP:
which bash
/usr/bin/bash
```

21. Check the details of the file system

```
/cow$ df -h
Filesystem      Size  Used Avail Use% Mounted on
none            3.9G  4.0K  3.9G  1% /mnt/wsl
drivers        379G  257G  122G  68% /usr/lib
/wsl/drivers   none      3.9G    0  3.9G  0% /usr/lib
modules        none      3.9G    0  3.9G  0% /usr/lib
/modules/5.15.153.1-microsoft-standard-WSL2
/dev/sdc       1007G   17G  939G  2% /
```

22. Check the lines, word count, and characters in a file using different options

```
/cow$ wc pink.txt
1 5 28 pink.txt
```

Result:

Successfully implemented and Tested Linux Commands.

Ex. No: 2	Install Hadoop
23rd July, 2024	

Aim:

To install Hadoop in Ubuntu.

Program:

Steps from this link : <https://medium.com/@abhikdey06/apache-hadoop-3-3-6-installation-on-ubuntu-22-04-14516bceec85>

1. Install Java:

- Check if Java is installed: `java -version`
- If Java is not installed, download and install it:

```

```
sudo apt update
sudo apt install openjdk-8-jdk
```

```

2. Download Hadoop:

- Go to the Apache Hadoop releases page and download the binary distribution (e.g., Hadoop 3.4.0):

```
wget https://downloads.apache.org/hadoop/common/hadoop-3.4.0/hadoop-3.4.0.tar.gz
```

- Extract the downloaded file:

```
tar -xzvf hadoop-3.4.0.tar.gz
```

3. Move Hadoop to the desired directory: `sudo mv hadoop-3.4.0 /opt/hadoop`

4. Set environment variables:

- Open the `~/.bashrc` file for editing: `nano ~/.bashrc`
- Add the following lines at the end of the file:

```
export HADOOP_HOME=/opt/hadoop  
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin
```

- Save and exit, then refresh your terminal: `source ~/.bashrc`

5. Configure Hadoop:

- Navigate to the Hadoop configuration directory: `cd $HADOOP_HOME/etc/hadoop/`
- Edit `hadoop-env.sh` to set the Java home path: `nano hadoop-env.sh`
- Add or modify the line to include your Java installation path: `export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64`

6. Format the NameNode: `hadoop namenode -format`

7. Start YARN: `start-yarn.sh`

8. Start all Hadoop services: `start-all.sh`

9. Check running processes: `jps`

10. Access Hadoop web interface: `http://localhost:9870`

11. Stop all Hadoop services when needed: `stop-all.sh`

Output:

```
hadoop@LAPTOP-MSOAV9AM:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [ LAPTOP-MSOAV9AM ]
Starting resourcemanager
Starting nodemanagers
hadoop@LAPTOP-MSOAV9AM:~$ jps
1169 DataNode
1874 NodeManager
1427 SecondaryNameNode
2326 Jps
1003 NameNode
1724 ResourceManager
```

Result:

Successfully Installed Hadoop on Ubuntu Linux.

Ex. No: 3	MapReduce in Hadoop
30th July, 2024	

MapReduce in Hadoop

Aim:

To implement MapReduce for the word count problem in Hadoop.

Program:

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
                       ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
                           Context context
                           ) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }
}
```

```

    }

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Program (Python):

Mapper.py

```
#!/usr/bin/env python3
```

```
import sys
```

```

# Input comes from standard input (stdin)
for line in sys.stdin:
    # Remove leading and trailing whitespace
    line = line.strip()
    # Split the line into words
    words = line.split()
    # Increase counters
    for word in words:
        # Write the results to standard output (stdout)
        # What we output here will be the input for the Reducer
        # The tab-delimited format is <word, 1>
        print('%s\t%s' % (word, 1))

```

reducer.py

```
#!/usr/bin/env python3
```

```
import sys

current_word = None
current_count = 0
word = None

# Input comes from standard input (stdin)
for line in sys.stdin:

    # Remove leading and trailing whitespace
    line = line.strip()

    # Parse the input we got from mapper.py
    word, count = line.split('\t', 1)

    # Convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # Count was not a number, so ignore/discard this line
        continue

    # This IF-switch only works because Hadoop sorts map output
    # by key (here: word) before it is passed to the reducer
    if current_word == word:
        current_count += count
    else:
        if current_word:
            # Write result to standard output (stdout)
            print('%s\t%s' % (current_word, current_count))
        current_count = count
        current_word = word

# Do not forget to output the last word if needed!
```

```
if current_word == word:  
    print('%s\t%s' % (current_word, current_count))
```

commands:

```
chmod +x mapper.py
```

```
chmod +x reducer.py
```

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/hadoop-streaming.jar \  
-input myInputDirs \  
-output myOutputDir \  
-mapper /bin/cat \  
-reducer /usr/bin/wc
```

```
mapred streaming -input /wordcount/input -output /wordcount/output4 -mapper  
.mapper.py -reducer ./reducer.py
```

Output:

```
hadoop@LAPTOP-MSOAV9AM:~$ hdfs dfs -cat /output2/part-00000  
1926      1  
a          1  
alexander      1  
am          2  
as          1  
great        2  
hello        1  
i            3  
in          1  
king         1  
my          1  
name         1  
suggests      1  
the          1  
was          1
```

Result:

Successfully implemented a MapReduce Function using Hadoop.

Ex. No: 4	
6th August, 2024	

MapReduce Problems

Aim:

1. To Implement map reduce for NCDC weather dataset using Hadoop and find the max and min temperature.
2. Implement Apriori algorithm using map reduce paradigm.

Program 1:

TemperatureMapper.java

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.FloatWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

import java.io.IOException;

public class TemperatureMapper extends Mapper<LongWritable, Text, Text, FloatWritable> {

    private final static FloatWritable temperature = new FloatWritable();
    private final Text tmaxKey = new Text("TMAX");
    private final Text tminKey = new Text("TMIN");

    @Override
    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();
        // Split the line on commas for CSV file
        String[] fields = line.split(",");
        if (fields.length > 1) {
            try {
                float temp = Float.parseFloat(fields[1]);
                if (temp > 0) {
                    if (temp > temperature.get()) {
                        temperature.set(temp);
                        context.write(tmaxKey, temperature);
                    }
                    if (temp < temperature.get()) {
                        temperature.set(temp);
                        context.write(tminKey, temperature);
                    }
                }
            } catch (NumberFormatException e) {
                System.out.println("Error: " + e.getMessage());
            }
        }
    }
}
```

```
// Skip header row and incomplete rows
if (fields[0].equals("STATION") || fields.length < 6) {
    return;
}

// Safely parse TMAX and TMIN values
try {
    String tmaxStr = fields[4].trim();
    tmaxStr = tmaxStr.replace("'", ' ');
    String tminStr = fields[5].trim();
    tminStr = tminStr.replace("'", ' ');
    // System.out.println(tmaxStr);

    if (!tmaxStr.isEmpty() && !tminStr.isEmpty()) {
        float tmax = Float.parseFloat(tmaxStr);
        float tmin = Float.parseFloat(tminStr);

        temperature.set(tmax);
        context.write(tmaxKey, temperature);
        temperature.set(tmin);
        context.write(tminKey, temperature);
    }
} catch (NumberFormatException e) {
    // Log parsing errors for debug purposes
    System.err.println("Error parsing temperature: " +
e.getMessage());
}
}

TemperatureReducer.java
```

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.FloatWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

import java.io.IOException;

public class TemperatureReducer extends Reducer<Text, FloatWritable,
Text, FloatWritable> {

    @Override
    protected void reduce(Text key, Iterable<FloatWritable> values,
Context context) throws IOException, InterruptedException {
        Float extremeValue;

        if (key.toString().equals("TMAX")) {
            extremeValue = Float.MIN_VALUE;
            for (FloatWritable value : values) {
                extremeValue = Math.max(extremeValue, value.get());
            }
            System.out.println("Maximum Temperature: " + extremeValue);
        } else { // TMIN
            extremeValue = Float.MAX_VALUE;
            for (FloatWritable value : values) {
                extremeValue = Math.min(extremeValue, value.get());
            }
            System.out.println("Minimum Temperature: " + extremeValue);
        }

        // Emit the result
        context.write(key, new FloatWritable(extremeValue));
    }
}
```

```
    }  
}  
}
```

TemperatureDriver.java

```
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.FloatWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
  
public class TemperatureDriver {  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Max and Min Temperature");  
  
        job.setJarByClass(TemperatureDriver.class);  
        job.setMapperClass(TemperatureMapper.class);  
        job.setReducerClass(TemperatureReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(FloatWritable.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Program 2 :

Mapper.py

```
#!/usr/bin/env python

import sys
from itertools import combinations

def generate_combinations(item_list, length):
    return list(combinations(item_list, length))

# Input comes from standard input (stdin)
for line in sys.stdin:
    line = line.strip()
    items = line.split()
    for length in range(1, len(items) + 1):
        for combination in generate_combinations(items, length):
            print(f"{''.join(combination)}\t1")
```

Reducer.py

```
#!/usr/bin/env python

import sys

current_itemset = None
current_count = 0

# Input comes from standard input (stdin)
for line in sys.stdin:
    line = line.strip()
    itemset, count = line.split('\t', 1)
```

```

count = int(count)

if current_itemset == itemset:
    current_count += count
else:
    if current_itemset:
        print(f"{current_itemset}\t{current_count}")
    current_count = count
    current_itemset = itemset

if current_itemset == itemset:
    print(f"{current_itemset}\t{current_count}")

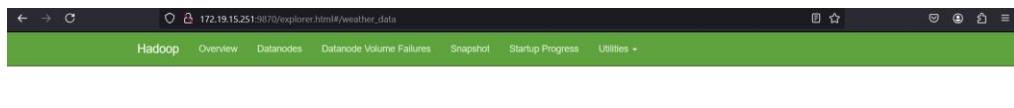
```

Outputs :

```

hadoop@LAPTOP-MSOAV9AM:~/hadoop/bin$ hdfs dfs -cat /weather_data/output/part-00000
Min Temperature: -11.17
Max Temperature: 29.67

```



Browse Directory								
<input type="text" value="/weather_data"/> <input type="button" value="Go!"/> <input type="button" value="New"/> <input type="button" value="Edit"/> <input type="button" value="Delete"/> 								
<input type="button" value="Search: "/>								
<input type="button" value="Show: 25 entries"/>								
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	-rw-r--r--	hadoop	supergroup	196.1 KB	Sep 08 18:45	1	128 MB	weather_BE.csv

Showing 1 to 1 of 1 entries

Previous Next

Hadoop, 2023.

Result:

Successfully implemented a Weather Analysis and Apriori Algorithm using MapReduce

Ex. No: 5	Apache Spark
13th August, 2024	

Aim:

To install Spark and PySpark and run the WordCount Program and Distribution of Ratings in MovieLens Dataset.

Program:

1. Download and Install Spark
 - a. wget <https://www.apache.org/dyn/closer.lua/spark/spark-3.5.0/spark-3.5.0-bin-hadoop3.tgz>
 - b. tar -xzf spark-3.5.0-bin-hadoop3.tgz
 - c. mv spark-3.5.0-bin-hadoop3 spark
2. Set Environment Variables
 - a. nano ~/.bashrc
 - b. export SPARK_HOME=~/spark
 - c. export PATH=\$PATH:\$SPARK_HOME/bin
 - d. export PYTHONPATH=\$SPARK_HOME/python:\$PYTHONPATH
 - e. export PYSPARK_PYTHON=python3
 - f. source ~/.bashrc
3. Running Programs : spark-submit file.py

WordCount Program

```
import findspark
findspark.init()

from pyspark.sql import SparkSession
from pyspark.conf import SparkConf

spark =
SparkSession.builder.master("local[*]").appName("WordCount").getOrCreate()
```

```
sc = spark.sparkContext

text = "./words.txt"

text_read = sc.textFile(text)

text_read.take(20)

words = text_read.flatMap(lambda x: x.split())

words.take(15)

words = words.map(lambda word: word.rstrip('.'))

words.take(15)

word_count = words.map(lambda word: (word, 1))
word_count.take(20)

word_count_red = word_count.reduceByKey(lambda x, y: (x+y)).sortByKey()

word_count_red = word_count_red.map(lambda x: (x[1], x[0]))

word_count_red.take(15)
```

MovieLens Dataset

```
import findspark
findspark.init()
```

```
from pyspark.sql import SparkSession
from pyspark.conf import SparkConf

spark =
SparkSession.builder.master("local[*]").appName("MovieLens").getOrCreate()

sc = spark.sparkContext

text = "./u.data"

text_read = spark.read.csv(text)

import pyspark.pandas as ps

data = ps.read_csv(text, sep="\t", names = ["userid", "movieid",
"rating", "idid"])

data

len(data.groupby(['userid']).describe()['rating'])
```

Output:

	userid	movieid	rating	idid
[(1, 'college'),	0	196	242	3 881250949
(1, 'i'),	1	186	302	3 891717742
(1, 'in'),	2	22	377	1 878887116
(1, 'is'),	3	244	51	2 880606923
(1, 'my'),	4	166	346	1 886397596
(2, 'nadar'),	5	298	474	4 884182806
(2, 'shiv'),	6	115	265	2 881171488
(1, 'study'),	7	253	465	5 891628467
(1, 'university')]	8	305	451	3 886324817
	9	6	86	3 883603013
	10	62	257	2 879372434
	11	286	1014	5 879781125
	12	200	222	5 876042340
	13	210	40	3 891035994
	14	224	29	3 888104457
	15	303	785	3 879485318
	16	122	387	5 879270459

Result:

Successfully Installed Spark and PySpark, and implemented the WordCount Program and Distribution of Movie Reviews using them.

Ex. No: 6	PySpark 2
20th August, 2024	

Aim:

- (a) To use the “friends _ test” dataset. Col1 is ID, Col2 is name, Col 3 is Age, Col 4 is num of friends. Understand mapvalues function of RDD in spark and find the average number of friends for each unique age present in the dataset.
- (b) Use the “temp.csv” dataset. Column headers are present in the dataset. Understand filter operations and filter out only the “TMIN” values from the “desc” column. With the resultant data (RDD) find the following:
 - a. Minimum temperature (overall)
 - b. Minimum temperature for every ItemID
 - c. Minimum temperature for every StationID.
- (c) Use the same dataset, filter only “TMAX” column and find the maximum temperatures just like the ones mentioned above.

Program:

```

from pyspark import SparkConf, SparkContext
from pyspark.sql import *
from pyspark.sql.functions import *
from pyspark.sql.types import *

spark = SparkSession.builder.appName("friends
test").config("spark.memory.offHeap.e
df = spark.read.csv('friends_test.csv',header=False)
df.explain()
spark.stop()
conf = SparkConf().setAppName("Basicapp").setMaster("local[*]")
sc = SparkContext(conf=conf)
rdd = sc.textFile("friends_test.csv")
rdd.first()
rdd_split = rdd.map(lambda line: line.split(","))
for row in rdd_split.take(5):
    print(row)
  
```

RDD approach for average number of friends

```
age_friends_rdd = rdd_split.map(lambda row: (int(row[2]), (int(row[3]), 1)))
sum_count_rdd = age_friends_rdd.reduceByKey(lambda a, b: (a[0] + b[0], a[1] + b[1]))
avg_friends_by_age = sum_count_rdd.mapValues(lambda x: x[0] / x[1])
sorted_avg_friends_by_age = avg_friends_by_age.sortByKey()
for age, avg_friends in sorted_avg_friends_by_age.collect():
    print(f"Age: {age}, Average Number of Friends: {avg_friends:.2f}")
age_friends_rdd = rdd_split.map(lambda row: (int(row[2]), int(row[3])))
max_friends_by_age = age_friends_rdd.reduceByKey(lambda a, b: a if a > b else b)
min_friends_by_age = age_friends_rdd.reduceByKey(lambda a, b: a if a < b else b)
max_min_friends_by_age = max_friends_by_age.join(min_friends_by_age)
sorted_max_min_friends_by_age = max_min_friends_by_age.sortByKey()
for age, (max_friends, min_friends) in sorted_max_min_friends_by_age.collect():
    print(f"Age: {age}, Max Friends: {max_friends}, Min Friends: {min_friends}")
sc.stop()
```

Temp dataset

```
conf = SparkConf().setAppName("temp_dataset").setMaster("local[*]")
sc = SparkContext(conf=conf)
rdd = sc.textFile("temp.csv")
rdd.first()
rdd_header = rdd.first()
rdd_filter = rdd.filter(lambda row: row != rdd_header)
rdd_data = rdd_filter.map(lambda row: row.split(","))
def rdd_display(x, threshold=5):
    count = 0
    for i in x.collect():
        print(i)
```

```

        count+=1
        if(count>threshold):
            break
    rdd_display(rdd_data)
    rdd_TMIN_filter = rdd_data.filter(lambda row:row[2]=="TMIN")
    rdd_display(rdd_TMIN_filter)

    rdd_min_overall = rdd_TMIN_filter.map(lambda x:int(x[3])).reduce(lambda
a,b:a if a
print("Minimum temperature overall",rdd_min_overall)
    rdd_min_itemID = rdd_TMIN_filter.map(lambda
x:(x[0],int(x[3]))).reduceByKey(lambda
print("minimum temperature by itemID")
    rdd_display(rdd_min_itemID)
    rdd_min_stationID = rdd_TMIN_filter.map(lambda
x:(x[1],int(x[3]))).reduceByKey(lambda
print("minimum temperature by StationID")
    rdd_display(rdd_min_stationID,10)
    rdd_TMAX_filter = rdd_data.filter(lambda row:row[2]=="TMAX")
    rdd_display(rdd_TMAX_filter)
    rdd_max_overall = rdd_TMAX_filter.map(lambda x:int(x[3])).reduce(lambda
a,b:a if a
print("Maximum temperature overall",rdd_max_overall)
    rdd_max_itemID = rdd_TMAX_filter.map(lambda
x:(x[0],int(x[3]))).reduceByKey(lambda
print("maximum temperature by itemID")
    rdd_display(rdd_max_itemID)
    rdd_max_stationID = rdd_TMAX_filter.map(lambda
x:(x[1],int(x[3]))).reduceByKey(lambda
print("maximum temperature by StationID")
    rdd_display(rdd_max_stationID,10)
sc.stop()

```

Output:

Age: 18, Average Number of Friends: 343.38
 Age: 19, Average Number of Friends: 213.27
 Age: 20, Average Number of Friends: 165.00
 Age: 21, Average Number of Friends: 350.88
 Age: 22, Average Number of Friends: 286.43
 Age: 23, Average Number of Friends: 246.30
 Age: 24, Average Number of Friends: 233.80
 Age: 25, Average Number of Friends: 197.45
 Age: 26, Average Number of Friends: 242.06
 Age: 27, Average Number of Friends: 228.12
 Age: 28, Average Number of Friends: 209.10
 Age: 29, Average Number of Friends: 215.92
 Age: 30, Average Number of Friends: 235.82
 Age: 31, Average Number of Friends: 267.25
 Age: 32, Average Number of Friends: 207.91
 Age: 33, Average Number of Friends: 325.33
 Age: 34, Average Number of Friends: 245.50
 Age: 35, Average Number of Friends: 211.62
 Age: 36, Average Number of Friends: 246.60
 Age: 37, Average Number of Friends: 249.33
 Age: 38, Average Number of Friends: 193.53
 Age: 39, Average Number of Friends: 169.29
 Age: 40, Average Number of Friends: 250.82
 Age: 41, Average Number of Friends: 268.56
 Age: 42, Average Number of Friends: 303.50
 Age: 43, Average Number of Friends: 230.57
 Age: 44, Average Number of Friends: 282.17
 Age: 45, Average Number of Friends: 309.54
 Age: 46, Average Number of Friends: 223.69
 Age: 47, Average Number of Friends: 233.22
 Age: 48, Average Number of Friends: 281.40
 Age: 49, Average Number of Friends: 184.67
 Age: 50, Average Number of Friends: 254.60
 Age: 51, Average Number of Friends: 302.14
 Age: 52, Average Number of Friends: 340.64
 Age: 53, Average Number of Friends: 222.86
 Age: 54, Average Number of Friends: 278.08
 Age: 55, Average Number of Friends: 295.54
 Age: 56, Average Number of Friends: 306.67
 Age: 57, Average Number of Friends: 258.83
 Age: 58, Average Number of Friends: 116.55
 Age: 59, Average Number of Friends: 220.00
 Age: 60, Average Number of Friends: 202.71
 Age: 61, Average Number of Friends: 256.22
 Age: 62, Average Number of Friends: 220.77
 Age: 63, Average Number of Friends: 384.00
 Age: 64, Average Number of Friends: 281.33
 Age: 65, Average Number of Friends: 298.20
 Age: 66, Average Number of Friends: 276.44
 Age: 67, Average Number of Friends: 214.62
 Age: 68, Average Number of Friends: 269.60
 Age: 69, Average Number of Friends: 235.20

Age: 18, Max Friends: 499, Min Friends: 24
 Age: 19, Max Friends: 404, Min Friends: 5
 Age: 20, Max Friends: 384, Min Friends: 1
 Age: 21, Max Friends: 491, Min Friends: 89
 Age: 22, Max Friends: 478, Min Friends: 6
 Age: 23, Max Friends: 392, Min Friends: 65
 Age: 24, Max Friends: 492, Min Friends: 49
 Age: 25, Max Friends: 485, Min Friends: 1
 Age: 26, Max Friends: 492, Min Friends: 2
 Age: 27, Max Friends: 471, Min Friends: 53
 Age: 28, Max Friends: 378, Min Friends: 32
 Age: 29, Max Friends: 367, Min Friends: 11
 Age: 30, Max Friends: 487, Min Friends: 17
 Age: 31, Max Friends: 481, Min Friends: 15
 Age: 32, Max Friends: 412, Min Friends: 24
 Age: 33, Max Friends: 471, Min Friends: 74
 Age: 34, Max Friends: 423, Min Friends: 48
 Age: 35, Max Friends: 428, Min Friends: 13
 Age: 36, Max Friends: 493, Min Friends: 49
 Age: 37, Max Friends: 471, Min Friends: 46
 Age: 38, Max Friends: 459, Min Friends: 2
 Age: 39, Max Friends: 275, Min Friends: 68
 Age: 40, Max Friends: 465, Min Friends: 7
 Age: 41, Max Friends: 397, Min Friends: 62
 Age: 42, Max Friends: 467, Min Friends: 95
 Age: 43, Max Friends: 428, Min Friends: 48
 Age: 44, Max Friends: 499, Min Friends: 61
 Age: 45, Max Friends: 497, Min Friends: 54
 Age: 46, Max Friends: 462, Min Friends: 63
 Age: 47, Max Friends: 488, Min Friends: 4
 Age: 48, Max Friends: 439, Min Friends: 57
 Age: 49, Max Friends: 476, Min Friends: 17
 Age: 50, Max Friends: 436, Min Friends: 119
 Age: 51, Max Friends: 493, Min Friends: 81
 Age: 52, Max Friends: 487, Min Friends: 77
 Age: 53, Max Friends: 451, Min Friends: 86
 Age: 54, Max Friends: 462, Min Friends: 7
 Age: 55, Max Friends: 474, Min Friends: 57
 Age: 56, Max Friends: 444, Min Friends: 15
 Age: 57, Max Friends: 465, Min Friends: 8
 Age: 58, Max Friends: 348, Min Friends: 6
 Age: 59, Max Friends: 439, Min Friends: 14
 Age: 60, Max Friends: 324, Min Friends: 2
 Age: 61, Max Friends: 469, Min Friends: 2
 Age: 62, Max Friends: 496, Min Friends: 12
 Age: 63, Max Friends: 469, Min Friends: 342
 Age: 64, Max Friends: 499, Min Friends: 65
 Age: 65, Max Friends: 443, Min Friends: 101
 Age: 66, Max Friends: 496, Min Friends: 41
 Age: 67, Max Friends: 445, Min Friends: 35
 Age: 68, Max Friends: 490, Min Friends: 21
 Age: 69, Max Friends: 491, Min Friends: 9

```

['ITE00100554', '18000101', 'TMAX', '-75']
['ITE00100554', '18000101', 'TMIN', '-148']
['GM000010962', '18000101', 'PRCP', '0']
['EZE00100082', '18000101', 'TMAX', '-86']
['EZE00100082', '18000101', 'TMIN', '-135']
['ITE00100554', '18000102', 'TMAX', '-60']

```

```

['ITE00100554', '18000101', 'TMIN', '-148']
['EZE00100082', '18000101', 'TMIN', '-135']
['ITE00100554', '18000102', 'TMIN', '-125']
['EZE00100082', '18000102', 'TMIN', '-130']
['ITE00100554', '18000103', 'TMIN', '-46']
['EZE00100082', '18000103', 'TMIN', '-73']

```

Minimum temperature overall -148

minimum temperature by itemID

('ITE00100554', -148)
 ('EZE00100082', -135)

```
minimum temperature by StationID
('18000102', -130)
('18000104', -74)
('18000106', -57)
('18000110', -75)
('18000111', -62)
('18000112', -60)
('18000114', -35)
('18000115', -23)
('18000116', -37)
('18000117', -35)
('18000118', 9)
```

```
['ITE00100554', '18000101', 'TMAX', '-75']
['EZE00100082', '18000101', 'TMAX', '-86']
['ITE00100554', '18000102', 'TMAX', '-60']
['EZE00100082', '18000102', 'TMAX', '-44']
['ITE00100554', '18000103', 'TMAX', '-23']
['EZE00100082', '18000103', 'TMAX', '-10']
```

```
maximum temperature by StationID
('18000102', -44)
('18000104', 0)
('18000106', 13)
('18000110', 46)
('18000111', 66)
('18000112', 41)
('18000114', 41)
('18000115', 54)
('18000116', 56)
('18000117', 84)
('18000118', 59)
```

Result:

Successfully understood map value functions and utilised them for various tasks

Ex. No: 7

27th August, 2024

Hadoop and Docker

Aim:

To install hadoop in docker containers

Output:

The terminal window shows the following content:

```
~/ex_docker
Dockerfile
dummy
hadoop-streaming-3.4.0.jar
mapper.py
reducer.py
wc.jar
WordCount.class
WordCount.java
WordCount$IntSumReducer.class
WordCount$TokenizerMapper.class
words.txt

SUBCOMMAND may print help when invoked w/o parameters or with -h.
hadoop@85dfacaf27ee:~/ex_docker$ hdfs dfs -cat /user/hadoop/output/part-r-00000
I          1
a          1
above     1
are        1
diamond   1
high       1
how        1
in         1
like       1
little     2
sky         1
so          1
star        2
the         2
twinkle    4
up          1
what        1
wonder     1
world      1
you        1
hadoop@85dfacaf27ee:~/ex_docker$
```

Name	Image	Status	Port(s)	CPU (%)	Last started	Actions
hadoop-container	hadoop image <none>	Running		3.16%	1 hour ago	[Container Actions]

Result:

Successfully installed hadoop in docker containers

Ex. No: 8	
3rd September, 2024	Public and Private Keys

Aim:

To secure an EC2 instance using SSH Keys and Network Access Control.

Program:

1. **Generate Private and Public Key Pairs**
 - o **Linux/Mac:**
 - Create a new directory for key pairs: `mkdir key-pair-labs && cd key-pair-labs`
 - Generate a private key: `openssl genrsa -out snu-privatekey.pem 2048`
 - Generate a public key from the private key: `openssl rsa -in snu-privatekey.pem -pubout -out snu-publickey.pem`
 - Set permissions on the private key: `chmod 400 snu-privatekey.pem`
 - Copy the public key: `cat snu-publickey.pem`
 - o Paste the public key into the AWS console.
2. **Launch an Ubuntu EC2 Instance:**
 - o Launch a new EC2 instance using the public key.
 - o Choose an appropriate instance type and AMI.
 - o Allocate an Elastic IP address and attach it to the instance.
3. **Login to the Instance:**
 - o Use the private key generated in step 1 to SSH into the instance.
 - o Linux / Mac : `ssh -i snu-privatekey.pem ubuntu@<public_ip_address>`
4. **Edit Security Group:**
 - o Open the Security Group settings for the instance.
 - o Add an inbound rule to allow ICMP traffic (ping) from anywhere:
 - Type: Custom TCP Rule
 - Protocol: ICMP
 - Port Range: All
 - Source: 0.0.0.0/0
5. **Ping the Instance:**
 - o Ping the public IP address of the instance to verify connectivity.

- You should be able to ping the instance successfully.

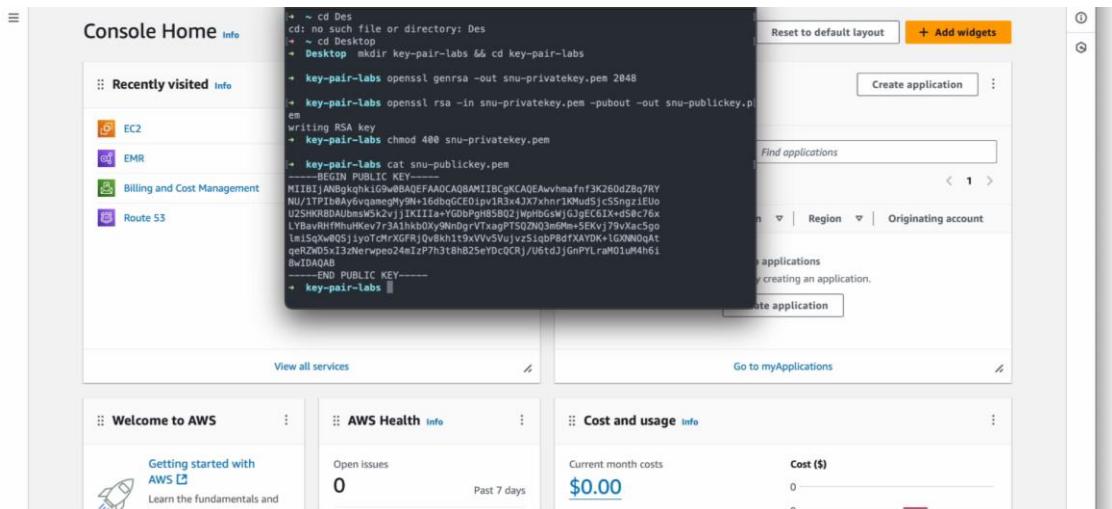
6. Edit Default NACL:

- Open the Network Access Control List (NACL) settings for the instance.
- Edit the default NACL to block ICMP traffic.
- Add an egress rule to deny ICMP traffic:
 - Type: Custom TCP Rule
 - Protocol: ICMP
 - Port Range: All
 - Destination: 0.0.0.0/0

7. Ping the Instance (Again):

- Ping the public IP address of the instance again.
- You should not be able to ping the instance this time, as ICMP traffic is now blocked.

Output:



The screenshot shows the AWS Management Console Home page. A terminal window is open in the top right corner, displaying the following command sequence:

```

+ ~ cd Des
cd: no such file or directory: Des
+ ~ cd Desktop
+ desktop mkdir key-pair-labs && cd key-pair-labs
+ key-pair-labs openssl genrsa -out snu-privatekey.pem 2048
+ key-pair-labs openssl rsa -in snu-privatekey.pem -pubout -out snu-publickey.pem
+ key-pair-labs chmod 400 snu-privatekey.pem
+ key-pair-labs cat snu-publickey.pem
-----BEGIN PUBLIC KEY-----
MIIBIjANBgkqhkiG9wBADEFAAQCA08AM1IBgkKCAQEAwhmafrf3K260dZBq7RY
NU/1TP1o@Ay6vgameghYn+16dqGCEOipvr1R3x4jX7xhmr1KMud5jc5ngzLEUo
U2SHKRB0AubmWSk2vjjIKIIia+YObDpgb$5B02jWptGswjGJEG6Ix+d5oC76x
LYBavRHfmuKeV7r3A1hkboKy+Nn0gVtXapPTSQZN03m6M+SEkjv79Xae5go
lmsQxW0QSj1yo7CMXGFrijv8kh19kvVv5VuUvzS1abP8dTXAYD+k+GxNNQdAt
qeH2mEiX3zNerwpe24m1zp7h3t8hB25eYdCQR/UTdJjGnPYLraM01uM4h6i
Bw1D0AQ
-----END PUBLIC KEY-----
+ key-pair-labs

```

The terminal window has a yellow border and is titled "Console Home". Below it, the main console interface shows sections for "Recently visited" services (EC2, EMR, Billing and Cost Management, Route 53), "Welcome to AWS" (Getting started with AWS), "AWS Health" (Open issues 0), and "Cost and usage" (Current month costs \$0.00).

Import settings

Name
key_pair_exercise
The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair file

Choose **Browse** and navigate to your public key. You may change the name of your key. Alternatively, paste the contents of your public key into the **Public key contents** text box.

```
MIIBJjANBgkqhkiG9w0BAQEAAQ8AMIBgkCgKCAQEAvvhmafnf5K26OdZ8q7RY
NU/1TPb0A0y6qamegM9h+16dbqGCCEOpvR1R3x4JXxmr1KMudj5SngzrEUo
U25HkR8DAUbmsW5k2vjjIKll+YGDbpH85BQ2jWpHtGswJGJgEC6X+d5Oc76x
LYBaRHfMuuhKev73A1hkbbOxyNnDpV7xapPTSQZnQ3mGm+5EKy/79Xac5go
ImISqkwQSjyoTCMrXGRjQvbkh19xVvVsUvjzSiqbP8dfXAYDX+IGXNNQoAt
qeRZWDSxi3zNerwpeo24mizP7h3t8hb25eYD;cQCRj/U6tdJjGnPVLraMO1uM4h6i
BwfDAQAB
```

Tags - optional
No tags associated with the resource.

You can add up to 50 more tags.

EC2 > Instances > i-0d9e44872c43f4bc0 (Key_Pair_exercise)

Instance summary for i-0d9e44872c43f4bc0 (Key_Pair_exercise) [Info](#)

Updated less than a minute ago

Instance ID i-0d9e44872c43f4bc0	Public IPv4 address 3.83.38.71 open address	Private IPv4 addresses 172.31.92.25
IPv6 address -	Instance state Running	Public IPv4 DNS ec2-3-83-38-71.compute-1.amazonaws.com open address
Hostname type IP name: ip-172-31-92-25.ec2.internal	Private IP DNS name (IPv4 only) ip-172-31-92-25.ec2.internal	Elastic IP addresses -
Answer private resource DNS name IPv4 (A)	Instance type t2.micro	AWS Compute Optimizer finding Opt-in to AWS Compute Optimizer for recommendations. Learn more
Auto-assigned IP address 3.83.38.71 [Public IP]	VPC ID vpc-0da85ce966b34be48	Auto Scaling Group name -
IAM Role -	Subnet ID subnet-06e2ed1d9731ecd56	
IMDSv2 Required	Instance ARN arn:aws:ec2:us-east-1:996326397401:instance/i-0d9e44872c43f4bc0	

[Details](#) [Status and alarms](#) [Monitoring](#) [Security](#) [Networking](#) [Storage](#) [Tags](#)

```

cd: ~ cd Desktop
Desktop/ ~ ssh -i sru-privkey.pem ubuntu@3.83.38.71
The authenticity of host '3.83.38.71 (3.83.38.71)' can't be established.
ED25519 key fingerprint is SHA256:bu@0R2s2X0cPMJSU5b5QfHb+s+n0VZyAhTuzxns8e57w.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '3.83.38.71' (ED25519) to the list of known hosts.

```

CloudShell Feedback

Instances (1/1)

Find Instance by Name: Instance ID: i-0d9e44872c43f4bc0

Name Key_Pair_exec

Instance summary

Instance ID: i-0d9e44872c43f4bc0

IPv6 address: -

Hostname type: IP name: ip-172-31-92-25.ec2.internal

Private IP DNS name (IPv4 only): ip-172-31-92-25.ec2.internal

Instance state: Running

Answer private resource DNS name: IPv4 (A)

Instance type: t2.micro

VPC ID: -

Elastic IP addresses: -

AWS Compute Optimizer finding: -

Private IPv4 addresses: 172.31.92.25

Public IPv4 DNS: ec2-3-83-38-71.compute-1.amazonaws.com | open address

CloudShell Feedback

Instances (1/1)

Find Instance by Name: Instance ID: i-0d9e44872c43f4bc0

Name Key_Pair_exec

Instance summary

Instance ID: i-0d9e44872c43f4bc0

IPv6 address: -

Hostname type: IP name: ip-172-31-92-25.ec2.internal

Private IP DNS name (IPv4 only): ip-172-31-92-25.ec2.internal

Instance state: Running

Answer private resource DNS name: IPv4 (A)

Instance type: t2.micro

VPC ID: -

Elastic IP addresses: -

AWS Compute Optimizer finding: -

Private IPv4 addresses: 172.31.92.25

Public IPv4 DNS: ec2-3-83-38-71.compute-1.amazonaws.com | open address

CloudShell Feedback

Edit inbound rules Info

Inbound rules control the incoming traffic that's allowed to reach the instance.

Inbound rules <small>Info</small>						
Security group rule ID	Type <small>Info</small>	Protocol <small>Info</small>	Port range	Source <small>Info</small>	Description - optional <small>Info</small>	
sgr-01f3f2dcce314d76e	SSH	TCP	22	Custom <input type="button" value="Delete"/>	<input type="text"/> 0.0.0.0/0 <input type="button" value="Delete"/>	<input type="button" value="Add rule"/>
-	All ICMP - IPv4	ICMP	All	Anyw... <input type="button" value="Delete"/>	<input type="text"/> 0.0.0.0/0 <input type="button" value="Delete"/>	<input type="button" value="Add rule"/>
-	All ICMP - IPv6	IPv6 ICMP	All	Anyw... <input type="button" value="Delete"/>	<input type="text"/> 0.0.0.0/0 <input type="button" value="Delete"/>	<input type="button" value="Add rule"/>

⚠ Rules with source of 0.0.0.0/0 or ::/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

CloudShell Feedback

Result:

Successfully Created Public and Private Keys and Utilised it to secure EC2 instances.

Ex. No: 9	
10th September, 2024	EC2

Aim:

To create an EC2 instance and Deploy a web server on it using Apache2

Program:

- 1. Create a Key Pair:**
 - o Use the AWS Management Console to create a new key pair.
 - o Download the private key file (.pem) and save it securely.
- 2. Launch EC2 Instance:**
 - o Ubuntu AMI,
 - o Choose an instance type based on your requirements.
 - o Configure security groups to allow SSH and HTTP traffic.
 - o Launch the instance and provide the private key file during the launch process.
- 3. Connect to Instance:**
 - o Use the SSH client to connect to the instance using the public IP address and private key file.
- 4. Update Package Lists:**
 - o Run the following command to update the package lists : sudo apt update
- 5. Install Apache2:** sudo apt install apache2
- 6. Check Apache Status:** sudo service apache2 status
- 7. Test Apache:**
 - o Open a web browser and enter the public IP address of the instance.
 - o You should see the default Apache2 welcome page.
- 8. Change the Contents of the Webpage:**
 - o cd /var/www/html/
 - o sudo chmod 777 index.html
 - o echo "website" > index.html
- 9. Test the website to show new content**

Output:

The screenshot shows the AWS EC2 Instances page. The instance summary for 'i-0ba1ed29c2be3ef84 (EC2_Lab)' is displayed. Key details include:

- Instance ID:** i-0ba1ed29c2be3ef84
- Public IPv4 address:** 184.73.151.28
- Instance state:** Running
- Private IPv4 addresses:** 172.31.93.220
- Public IPv4 DNS:** ec2-184-73-151-28.compute-1.amazonaws.com
- Hostname type:** IP name: ip-172-31-93-220.ec2.internal
- Private IP DNS name (IPv4 only):** ip-172-31-93-220.ec2.internal
- Instance type:** t2.micro
- VPC ID:** vpc-0da85ce966b34be48
- Elastic IP addresses:** -
- AWS Compute Optimizer finding:** Opt-in to AWS Compute Optimizer for recommendations.
- Auto Scaling Group name:** -
- IAM Role:** -
- Subnet ID:** subnet-06e2ed1d9731ecd56
- Instance ARN:** arn:aws:ec2:us-east-1:996326397401:instance/i-0ba1ed29c2be3ef84

```

→ Desktop mkdir key-pair-labs && cd key-pair-labs

→ key-pair-labs openssl genrsa -out snu-privatekey.pem 2048

[→ key-pair-labs openssl rsa -in snu-privatekey.pem -pubout -out snu-publickey.pem
em
writing RSA key
→ key-pair-labs chmod 400 snu-privatekey.pem

[→ key-pair-labs cat snu-publickey.pem
-----BEGIN PUBLIC KEY-----
MIIBIjANBggkqhkiG9w0BAQEFAAOCAQ8AMIBCgKCAQEAn5V9u25EVk5nNfs0qMoi
shWlxwjp9A1MGYOHIQjQdInUdk1t4U1NR57m+9YlRwaYyYojimEn1IJcVdHGzTl
R3BbGmPSKhR7ilawYKwb3IoCP3s1NPbg/Mq88PjzEdKPPjUP9xisnnNCkLNkoTwo
tnpXEtWUXdJ/8pWwwaGVQsUWHjN8YHbL3/oRoVndhWps9BGn4cqRV1DMpq5QwkKI
+iApDy4sj3xg24fha0pnbgZDYJSikarqLvT4TVISuNqSuxwj0qoMGZhkmold8Tcy
I873xrMgcAmmxwRfQ3FAhD39dRroatZPNvhDL0TmjEFVee1RpVrs5QDyFCRVa08Y
uwIDAQAB
-----END PUBLIC KEY-----
[→ key-pair-labs ssh -i snu-privatekey.pem ubuntu@184.73.151.28
The authenticity of host '184.73.151.28 (184.73.151.28)' can't be established.
ED25519 key fingerprint is SHA256:0447y6SkAs3URJwfWj4XN8W9P1GDvUeMPMA04QdgfnY.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
]

```

```
[ubuntu@ip-172-31-93-220:~$ sudo apt update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [12
6 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease [
126 kB]
Get:4 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:5 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Packag
es [15.0 MB]
Get:6 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe Translation-
en [5982 kB]
Get:7 http://security.ubuntu.com/ubuntu noble-security/main amd64 Packages [433
kB]
Get:8 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Compon
ents [3871 kB]
Get:9 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 c-n-f
Metadata [301 kB]
Get:10 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse amd64 Pac
kages [269 kB]
Get:11 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse Translati
on-en [118 kB]
Get:12 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse amd64 Com
ponents [35.0 kB]
Get:13 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse amd64 c-n
```

```
[ubuntu@ip-172-31-93-220:~$ sudo apt install apache2
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  apache2-bin apache2-data apache2-utils libapr1t64 libaprutil1-dbd-sqlite3
  libaprutil1-ldap libaprutil1t64 liblua5.4-0 ssl-cert
Suggested packages:
  apache2-doc apache2-suexec-pristine | apache2-suexec-custom www-browser
The following NEW packages will be installed:
  apache2 apache2-bin apache2-data apache2-utils libapr1t64
  libaprutil1-dbd-sqlite3 libaprutil1-ldap libaprutil1t64 liblua5.4-0 ssl-cert
0 upgraded, 10 newly installed, 0 to remove and 28 not upgraded.
Need to get 2084 kB of archives.
After this operation, 8094 kB of additional disk space will be used.
[Do you want to continue? [Y/n] Y
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 li
  bapr1t64 amd64 1.7.2-3.1ubuntu0.1 [108 kB]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 libaprutil
  1t64 amd64 1.6.3-1.1ubuntu7 [91.9 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 libaprutil
  1-dbd-sqlite3 amd64 1.6.3-1.1ubuntu7 [11.2 kB]
Get:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 libaprutil
  1-ldap amd64 1.6.3-1.1ubuntu7 [9116 B]
```

```
No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
[ubuntu@ip-172-31-93-220:~$ sudo service apache2 status
● apache2.service - The Apache HTTP Server
  Loaded: loaded (/usr/lib/systemd/system/apache2.service; enabled; preset: >
  Active: active (running) since Sun 2024-10-27 07:58:20 UTC; 14s ago
    Docs: https://httpd.apache.org/docs/2.4/
  Main PID: 2060 (apache2)
    Tasks: 55 (limit: 1130)
   Memory: 5.4M (peak: 5.7M)
      CPU: 32ms
     CGroup: /system.slice/apache2.service
             └─2060 /usr/sbin/apache2 -k start
                  ├─2063 /usr/sbin/apache2 -k start
                  ├─2064 /usr/sbin/apache2 -k start

Oct 27 07:58:20 ip-172-31-93-220 systemd[1]: Starting apache2.service - The Apa>
Oct 27 07:58:20 ip-172-31-93-220 systemd[1]: Started apache2.service - The Apac>
[lines 1-15/15 (END)]
```



21 of these updates are standard security updates.
To see these additional updates run: apt list --upgradable

Enable ESM Apps to receive additional future security updates.
See <https://ubuntu.com/esm> or run: sudo pro status

```
Last login: Sun Oct 27 07:57:17 2024 from 49.43.251.85
[ubuntu@ip-172-31-93-220:~$ cd /var/www/html/
[ubuntu@ip-172-31-93-220:/var/www/html$ sudo chmod 777 index.html
Command 'suo' not found, did you mean:
  command 'su' from deb util-linux (2.39.3-9ubuntu6.1)
  command 'sur' from deb subtle (0.11.3224-xi-2.2build5)
  command 'sup' from deb sup (20100519-3)
  command 'su1' from deb hxtools (20231101-1)
  command 'sudo' from deb sudo (1.9.14p2-1ubuntu1)
  command 'sudo' from deb sudo-ldap (1.9.14p2-1ubuntu1)
  command 'sumo' from deb sumo (1.18.0+dfsg-3build2)
  command 'zuo' from deb zuo (1.9-1)
  command 'sum' from deb coreutils (9.4-2ubuntu2)
Try: sudo apt install <deb name>
[ubuntu@ip-172-31-93-220:/var/www/html$ sudo chmod 777 index.html
[ubuntu@ip-172-31-93-220:/var/www/html$ echo "website" > index.html
ubuntu@ip-172-31-93-220:/var/www/html$
```



Result:

Successfully deployed an EC2 instance and tested using Apache Webservices.

Ex. No: 10	Route53
1st October, 2024	

Aim:

To set up a web server on an AWS EC2 instance and configure a domain name using GoDaddy and Route 53

Program:

1. **Login to GoDaddy and Create a subdomain:**
 - o Use your 12-digit registration number.
 - o Example: 21100101001.ngaws.xyz
2. **Create a Hosted Zone in Route 53:**
 - o Navigate to the Route 53 service in the AWS console.
 - o Create a new Hosted Zone for your subdomain.
3. **Get the Name Server information:**
 - o Retrieve the name server information from the Route 53 dashboard.
 - o Example: ns-1234.awsdns-12.org, [invalid URL removed]
4. **Update the NS record in GoDaddy:**
 - o Log in to the GoDaddy portal.
 - o Navigate to the DNS settings for your subdomain.
 - o Update the NS records with the name servers obtained from Route 53.
5. **Create an EC2 instance:**
 - o Launch an EC2 instance with an elastic public IP address.
6. **Install Apache:**
 - o Connect to the EC2 instance using SSH.
 - o Update the package lists: sudo apt update
 - o Install Apache: sudo apt install apache2
7. **Configure Apache:**
 - o Edit the Apache configuration file: sudo nano /etc/apache2/sites-available/000-default.conf
 - o Modify the DocumentRoot and ServerName directives to point to your desired web content directory and domain name.
 - o Save the configuration file and restart Apache: sudo systemctl restart apache2

8. Create an A record in Route 53:

- o Navigate to your Hosted Zone in Route 53.
- o Create a new A record with the following details:
 - Name: @ (for the root domain)
 - Value: The public IP address of your EC2 instance
 - TTL: 3600 (1 hour)

Output:

The screenshot shows the AWS Route 53 service dashboard. On the left, a sidebar lists various options like Dashboard, Hosted zones, Health checks, Profiles, IP-based routing, Traffic flow, Domains, Resolver, VPCs, and Outposts. The 'Hosted zones' section is selected. In the main pane, a success message at the top states: "21011101055.ngaws.xyz was successfully created. Now you can create records in the hosted zone to specify how you want Route 53 to route traffic for your domain." Below this, the '21011101055.ngaws.xyz' hosted zone is shown. The 'Records' tab is selected, displaying two entries:

Type	Name	Value	TTL
NS	21011101...	ns-124.awsdns-15.com. ns-999.awsdns-60.net. ns-2045.awsdns-63.co.uk. ns-1475.awsdns-56.org.	600 seconds
SOA	21011101...	ns-124.awsdns-15.com. awsd...	600 seconds

The screenshot shows the Cloudflare DNS interface. The left sidebar includes 'Domains', 'Portfolio', 'DNS', 'Transfers', 'Services', 'Tools', and 'Settings'. The 'DNS' section is active, showing a table of records. A success message box is overlaid on the right side of the table, stating: "Your DNS record has been updated successfully. Most DNS updates take effect within an hour, but could take up to 48 hours to update globally." The table contains the following records:

Type	Name	Value	TTL	Action
A	jenkins	13.233.133.184	600 seconds	Can't edit
A	varun-todo	3.108.220.63	600 seconds	Can't edit
NS	@	ns51.domaincontrol.com.	1 Hour	Can't delete
NS	@	ns52.domaincontrol.com.	1 Hour	Can't delete
NS	21011101055	ns-124.awsdns-15.com.	1 Hour	Delete Edit
NS	21011101074	ns-116.awsdns-14.com.	1 Hour	Delete Edit
NS	21011101075	ns-127.awsdns-15.com.	1 Hour	Delete Edit
NS	21011101079	ns-289.awsdns-36.com.	1 Hour	Delete Edit
NS	21011101079	ns-510.awsdns-63.com.	1 Hour	Delete Edit

Screenshot of the AWS EC2 Instances page showing the details for instance i-09a90d79acc628dbd.

Instance summary for i-09a90d79acc628dbd

Updated less than a minute ago

Attribute	Value
Instance ID	i-09a90d79acc628dbd
IPv6 address	-
Hostname type	IP name: ip-172-31-45-111.ec2.internal
Answer private resource DNS name	IPv4 (A)
Auto-assigned IP address	54.162.240.174 [Public IP]
IAM Role	-
IMDsv2 Required	Subnet ID subnet-08c80c8826320fef
VPC ID	vpc-0da85ce966b34be48
Instance state	Running
Private IP DNS name (IPv4 only)	ip-172-31-45-111.ec2.internal
Instance type	t2.micro
Private IPv4 address	54.162.240.174 open address
Public IPv4 address	172.31.45.111
Public IPv4 DNS	ec2-54-162-240-174.compute-1.amazonaws.com open address
Elastic IP addresses	-
AWS Compute Optimizer finding	Opt-in to AWS Compute Optimizer for recommendations.
Auto Scaling Group name	-
Instance ARN	arn:aws:ec2:us-east-1:1996326397401:instance/i-09a90d79acc628dbd

Details | Status and alarms | Monitoring | Security | Networking | Storage | Tags

Screenshot of the AWS Route 53 Hosted Zones page showing the records for the zone 21011101055.ngaws.xyz.

Route 53

Record for 21011101055.ngaws.xyz was successfully created.

Route 53 propagates your changes to all of the Route 53 authoritative DNS servers within 60 seconds. Use "View status" button to check propagation status.

Hosted zone details

Records (3) Info

Automatic mode is the current search behavior optimized for best filter results. To change modes go to settings.

Record	Type	Routing policy	Alias	Value/Route traffic to
21011101...	NS	Simple	No	ns-124.awsdns-15.com. ns-999.awsdns-60.net. ns-2045.awsdns-63.co.uk. ns-1475.awsdns-56.org.
21011101...	SOA	Simple	No	ns-124.awsdns-15.com. awsd...
www.210...	A	Simple	No	54.162.240.174

```
→ ~ links www.21011101055.ngaws.xyz
zsh: command not found: links
→ ~ nslookup www.21011101055.ngaws.xyz
Server:          172.16.0.2
Address:         172.16.0.2#53

Non-authoritative answer:
Name:   www.21011101055.ngaws.xyz
Address: 54.162.240.174
```

Result:

Successfully implemented a Route53 DNS Lookup on AWS.

Ex. No: 11	IAM
8th October, 2024	

Aim:

To implement IAM user access in AWS Console.

Program:

1. Create User Groups server_admin and dns_admin. Give them full access to EC2 Service and Route53 services.
2. Create users, alice, bob, cathay and david. Set passwords for them
3. Add alice and bob to user group server_adin and cathy and david to user group dns_admin.
4. Create user eve and give him Billing access.
5. Create user vishwa and give full (admin) access to the services.
6. Create an alias name for your account.
7. Use the alias URL to login to your account instead of the account ID.

Output:

Screenshot of the AWS IAM User Groups creation interface.

User group name: dns_admin

Add users to the group - Optional (0) Info:

Attach permissions policies - Optional (1/959) Info:

- route53fu (AWS managed) - Provides full access to all Amazon Route 53 resources.

Create user group button is visible.

Screenshot of the AWS IAM User creation review step.

User details: User name: alice, Console password type: Autogenerated, Require password reset: Yes.

Permissions summary:

- iamUserChangePassword (AWS managed) - Permissions policy
- server_admin (Group) - Permissions group

Tags - optional: No tags associated with the resource.

Add new tag button is visible.

Screenshot of the AWS IAM Users list after successful creation.

User created successfully: You can view and download the user's password and email instructions for signing in to the AWS Management Console.

Users (6) Info:

User Name	Path	Group	Last Activity	MFA	Password Age	Console Last Sign-in
alice	/	1	-	-	-	-
bob	/	1	-	-	-	-
cathy	/	1	-	-	-	-
david	/	1	-	-	-	-
eve	/	0	-	-	-	-

You are currently using the improved sign in UI experience.
The Improved sign in [?] experience will launch soon. During this time, you can still change back to legacy sign in using the dropdown in the upper right corner.

New sign in ▾ English ▾

aws

IAM user sign in ⓘ

Account ID (12 digits) or account alias

IAM username

Password

Show Password [Having trouble?](#)

Sign in

[Sign in using root user email](#)

[Create a new AWS account](#)

Remember this account

Amazon Lightsail
Lightsail is the easiest way to get started on AWS

[Learn more »](#)

aws

You must change your password to continue

AWS account 996326397401

IAM user name alice

Old password

New password

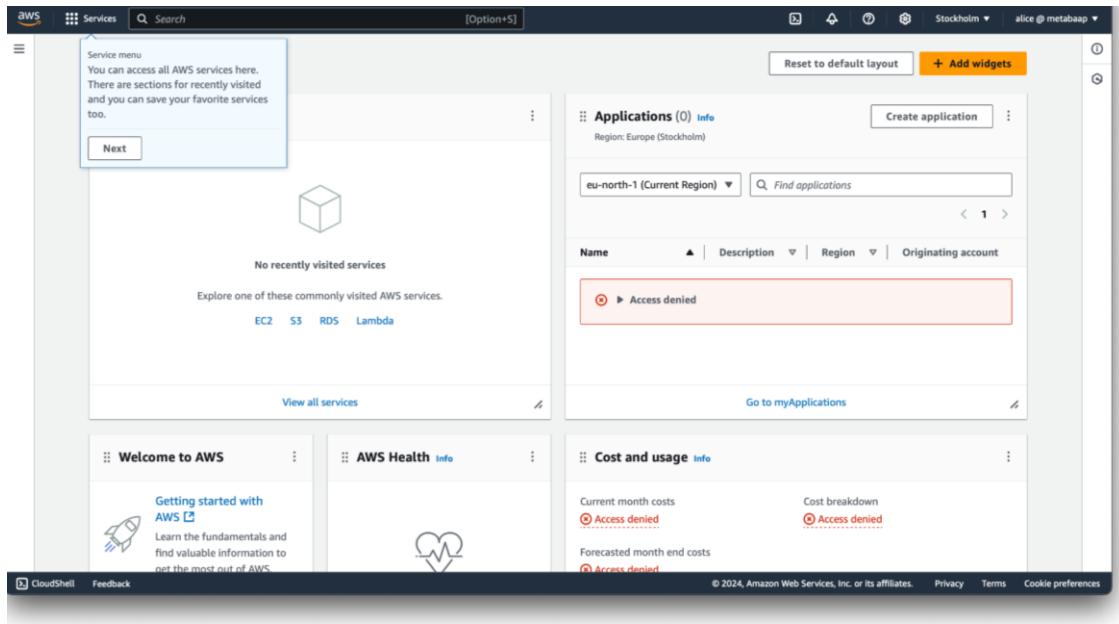
Retype new password

Confirm password change

[Sign in using root user email](#)

English ▾

Terms of Use Privacy Policy © 1996-2024, Amazon Web Services, Inc. or its affiliates.



Result:

Successfully implemented IAM labs in AWS